



# NVIDIA VIDEO CODEC SDK - DECODER

vNVDECODEAPI\_PG-08085-001\_v07 | June 2020

## Programming Guide



# TABLE OF CONTENTS

<b>Chapter 1. Overview.....</b>	<b>1</b>
1.1. Supported Codecs.....	1
<b>Chapter 2. Video Decoder Capabilities.....</b>	<b>3</b>
<b>Chapter 3. Video Decoder Pipeline.....</b>	<b>5</b>
<b>Chapter 4. Using NVIDIA Video Decoder (NVDECODE API).....</b>	<b>7</b>
4.1. Video Parser.....	7
4.1.1. Creating a parser.....	7
4.1.2. Parsing the packets.....	8
4.1.3. Destroying parser.....	8
4.2. Video Decoder.....	9
4.2.1. Querying decode capabilities.....	9
4.2.2. Creating a Decoder.....	10
4.2.3. Decoding the frame/field.....	12
4.2.4. Preparing the decoded frame for further processing.....	12
4.2.5. Querying the decoding status.....	14
4.2.6. Reconfiguring the decoder.....	14
4.2.7. Destroying the decoder.....	15
4.3. Writing an Efficient Decode Application.....	15

# Chapter 1.

## OVERVIEW

NVIDIA GPUs - beginning with the NVIDIA® Fermi™ generation - contain a video decoder engine (referred to as NVDEC in this document) which provides fully-accelerated hardware video decoding capability. NVDEC can be used for decoding bitstreams of various formats: H.264, HEVC (H.265), VP8, VP9, MPEG-1, MPEG-2, MPEG-4 and VC-1. NVDEC runs completely independent of compute/graphics engine.

NVIDIA provides software API and libraries for programming NVDEC. The software API, hereafter referred to as NVDECODER API lets developers access the video decoding features of NVDEC and interoperate NVDEC with other engines on the GPU.

NVDEC decodes the compressed video streams and copies the resulting YUV frames to video memory. With frames in video memory, video post processing can be done using CUDA. The NVDECODER API also provides CUDA-optimized implementation of commonly used post-processing operations such as scaling, cropping, aspect ratio conversion, de-interlacing and color space conversion to many popular output video formats. The client can choose to use the CUDA-optimized implementations provided by the NVDECODER API for these post-processing steps or choose to implement their own post-processing on the decoded output frames.

Decoded video frames can be presented to the display with graphics interoperability for video playback, passed directly to a dedicated hardware encoder (NVENC) for high-performance video transcoding, used for GPU accelerated inferencing or consumed further by CUDA or CPU-based processing.

## 1.1. Supported Codecs

The codecs supported by NVDECODER API are:

- ▶ MPEG-1,
- ▶ MPEG-2,
- ▶ MPEG4,
- ▶ VC-1,
- ▶ H.264 (AVCHD) (8 bit),
- ▶ H.265 (HEVC) (8bit, 10 bit and 12 bit),

- ▶ VP8,
- ▶ VP9(8bit, 10 bit and 12 bit).
- ▶ Hybrid (CUDA + CPU) JPEG

Refer to Chapter 2 for complete details about the video capabilities for various GPUs.

# Chapter 2.

## VIDEO DECODER CAPABILITIES

Table 1 shows the codec support and capabilities of the hardware video decoder for each GPU architecture.

Table 1 Hardware Video Decoder Capabilities>

GPU Architecture	MPEG-1 & MPEG-2	VC-1 & MPEG-4	H.264/AVCHD	H.265/HEVC	VP8	VP9
Fermi (GF1xx)	Maximum Resolution: 4080x4080	Maximum Resolution: 2048x1024 & 1024x2048	Maximum Resolution: 4096x4096 Profile: Baseline, Main, High profile up to Level 4.1	Unsupported	Unsupported	Unsupported
Kepler (GK1xx)	Maximum Resolution: 4080x4080	Maximum Resolution: 2048x1024 & 1024x2048	Maximum Resolution: 4096x4096 Profile: Main, Highprofile up to Level4.1	Unsupported	Unsupported	Unsupported
First generation Maxwell (GM10x)	Maximum Resolution: 4080x4080	Maximum Resolution: 2048x1024 & 1024x2048	Maximum Resolution: 4096x4096 Profile: Baseline, Main, High profile up to Level5.1	Unsupported	Unsupported	Unsupported
Second generation Maxwell (GM20x, except GM206)	Maximum Resolution: 4080x4080	Maximum Resolution: 2048x1024 & 1024x2048 Max bitrate: 60 Mbps	Maximum Resolution: 4096x4096 Profile: Baseline, Main, High profile up to Level5.1	Unsupported	Maximum Resolution: 4096x4096	Unsupported
GM206	Maximum Resolution: 4080x4080	Maximum Resolution: 2048x1024 & 1024x2048	Maximum Resolution: 4096x4096 Profile: Baseline, Main, High profile up to Level5.1	Maximum Resolution: 4096x2304 Profile: Main profile up to Level5.1 and main10 profile	Maximum Resolution: 4096x4096	Maximum Resolution: 4096x2304 Profile: Profile 0
GP100	Maximum Resolution: 4080x4080	Maximum Resolution: 2048x1024 & 1024x2048	Maximum Resolution: 4096x4096 Profile: Baseline, Main, High profile up to Level 5.1	Maximum Resolution: 4096x4096 Profile: Main profile up to Level 5.1, main10 and main12 profile	Maximum Resolution: 4096x4096	Maximum Resolution: 4096x4096 Profile: Profile 0

GPU Architecture	MPEG-1 & MPEG-2	VC-1 & MPEG-4	H.264/AVCHD	H.265/HEVC	VP8	VP9
GP10x/GV100/ Turing/GA100	Maximum Resolution: 4080x4080	Maximum Resolution: 2048x1024 & 1024x2048	Maximum Resolution: 4096x4096 Profile: Baseline, Main, High profile up to Level 5.1	Maximum Resolution: 8192x8192 Profile: Main profile up to Level 5.1, main10 and main12 profile	Maximum Resolution: 4096x4096 <sup>[1]</sup>	Maximum Resolution: 8192x8192 <sup>[2]</sup> Profile: Profile 0, 10-bit and 12-bit decoding

[1] Supported only on GP104, Turing and GA100

[2] VP9 10-bit and 12-bit decoding is supported on select GP10x GPUs

## Chapter 3.

# VIDEO DECODER PIPELINE

Decoder pipeline consists of three major components - Demuxer, Video Parser, and Video Decoder. The components are not dependent on each other and hence can be used independently. NVDECODE API provide API's for NVIDIA video parser and NVIDIA video decoder. Of these, NVIDIA video parser is purely a software component and user can plug-in any parser in place of NVIDIA video parser, if required (e.g. FFmpeg parser).

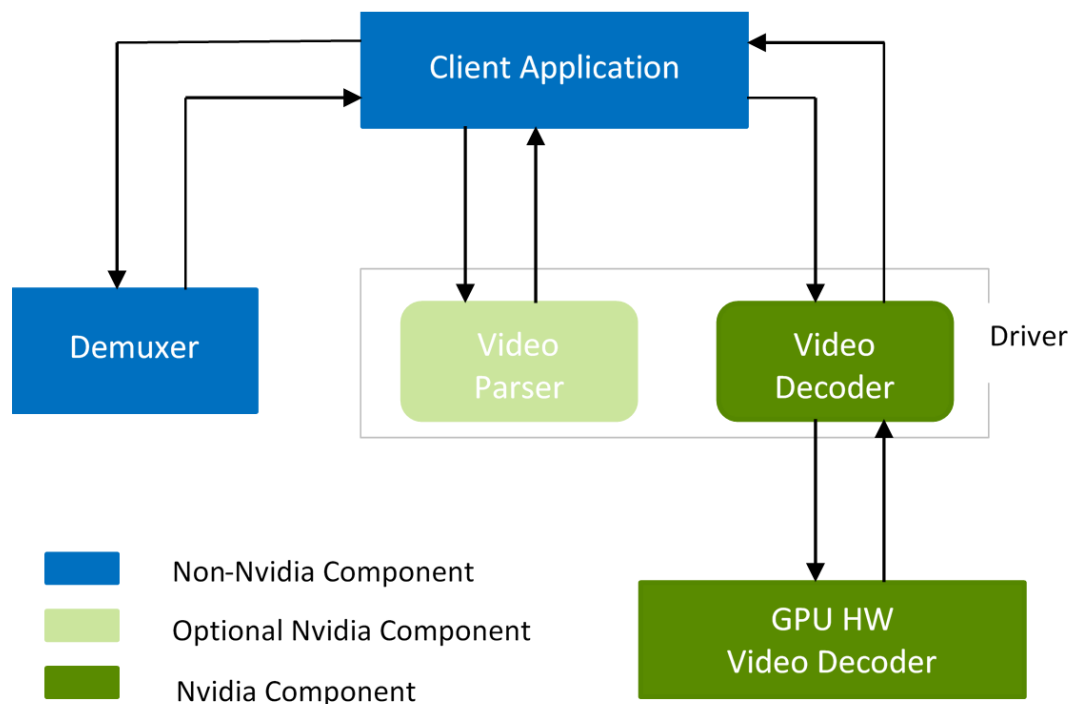


Figure 1 Video decoder pipeline using NVDECODE API

At a high level the following steps should be followed for decoding any video content using NVDECODEAPI:

1. Create a CUDA context.
2. Query the decode capabilities of the hardware decoder.

3. Create the decoder instance(s).
4. De-Mux the content (like .mp4). This can be done using third party software like FFMPEG.
5. Parse the video bitstream using parser provided by NVDECODE API or third-party parser such as FFmpeg.
6. Kick off the Decoding using NVDECODE API.
7. Obtain the decoded YUV for further processing.
8. Query the status of the decoded frame.
9. Depending on the decoding status, use the decoded output for further processing like rendering, inferencing, postprocessing etc.
10. If the application needs to display the output,
  - ▶ Convert decoded YUV surface to RGBA.
  - ▶ Map RGBA surface to DirectX or OpenGL texture.
  - ▶ Draw texture to screen.
11. Destroy the decoder instance(s) after the completion of decoding process.
12. Destroy the CUDA context.

The above steps are explained in the rest of the document and demonstrated in the sample application(s) included in the Video Codec SDK package.



# Chapter 4.

## USING NVIDIA VIDEO DECODER (NVDECODE API)

All NVDECODE APIs are exposed in two header-files: **cuviddec.h** and **nvcuvid.h**. These headers can be found under `..\Samples\NvCodec\NvDecoder` folder in the Video Codec SDK package. The samples in NVIDIA Video Codec SDK statically load the library (which ships as a part of the SDK package for windows) functions and include **cuviddec.h** and **nvcuvid.h** in the source files. The Windows DLL **nvcuvid.dll** is included in the NVIDIA display driver for Windows. The Linux library **libnvcuvid.so** is included with NVIDIA display driver for Linux.

The following sections in this chapter explain the flow that should be followed to accelerate decoding using NVDECODE API.

### 4.1. Video Parser

#### 4.1.1. Creating a parser

Parser object can be created by calling **cuvidCreateVideoParser()** after filling the structure **CUVIDPARSERPARAMS**. The structure should be filled up with following information about the stream to be decoded:

- ▶ **CodecType**: must be from **enum cudaVideoCodec**, indicating codec type of content like H.264, HEVC, VP9 etc.
- ▶ **ulMaxNumDecodeSurfaces**: This is number of surfaces in parser's DPB (decode picture buffer). This value may not be known at the parser initialization time and can be set to a dummy number like 1 to create parser object. Application must register a callback **pfnSequenceCallback** with the driver, which is called by the parser when the parser encounters the first sequence header or any changes in the sequence. This callback reports the minimum number of surfaces needed by parser's DPB for correct decoding in **CUVIDEOFORMAT::min\_num\_decode\_surfaces**. The sequence callback may return this value to the parser if wants to update **CUVIDPARSERPARAMS::ulMaxNumDecodeSurfaces**. The parser then overwrites **CUVIDPARSERPARAMS::ulMaxNumDecodeSurfaces** with the value

returned by the sequence callback, if return value of the sequence callback is greater than 1 (see description about **pfnSequenceCallback** below). Therefore, for optimum memory allocation, decoder object creation should be deferred until **CUVIDPARSERPARAMS::ulMaxNumDecodeSurfaces** is known, so that the decoder object can be created with required number of buffers, such that **CUVIDDECODERCREATEINFO::ulNumDecodeSurfaces = CUVIDPARSERPARAMS::ulMaxNumDecodeSurfaces**.

- ▶ **ulMaxDisplayDelay**: Max display callback delay. 0 = no delay
- ▶ **pfnSequenceCallback**: Application must register a function to handle any sequence change. Parser triggers this callback for initial sequence header or when it encounters a video format change. Return value from sequence callback is interpreted by the driver as follows:
  - ▶ 0: fail
  - ▶ 1: succeeded, but driver should not override **CUVIDPARSERPARAMS::ulMaxNumDecodeSurfaces**
  - ▶ >1: succeeded, and driver should override **CUVIDPARSERPARAMS::ulMaxNumDecodeSurfaces** with this return value
- ▶ **pfnDecodePicture**: Parser triggers this callback when bitstream data for one frame is ready. In case of field pictures, there may be two decode calls per one display call since two fields make up one frame. Return value from this callback is interpreted as:
  - ▶ 0: fail
  - ▶ ≥1: succeeded
- ▶ **pfnDisplayPicture**: Parser triggers this callback when a frame in display order is ready. Return value from this callback is interpreted as:
  - ▶ 0: fail
  - ▶ ≥1: succeeded

## 4.1.2. Parsing the packets

Packets extracted from demultiplexer are fed into parser using **cuvidParseVideoData()**. Parser triggers callbacks registered while creating parser object synchronously from within **cuvidParseVideoData()**, whenever there is sequence change or a picture is ready to be decoded and/or displayed. If the callback returns failure, it will be propagated by **cuvidParseVideoData()** to the application.

The decoded result gets associated with a picture-index value in the **CUVIDPICPARAMS** structure, which is also provided by the parser. This picture index is later used to map the decoded frames to CUDA memory.

## 4.1.3. Destroying parser

The user needs to call **cuvidDestroyVideoParser()** to destroy the parser object and free up all the allocated resources.

## 4.2. Video Decoder

### 4.2.1. Querying decode capabilities

The API `cuvidGetDecoderCaps()` lets users query the capabilities of underlying hardware video decoder.

As illustrated in Table 1, different GPUs have hardware decoders with different capabilities. Therefore, to ensure your application works on all generations of GPU hardware, it is highly recommended that the application queries the hardware capabilities and makes appropriate decision based on presence/absence of the desired capability/functionality.

The API `cuvidGetDecoderCaps()` lets users query the capabilities of underlying hardware video decoder. Calling thread should have a valid CUDA context associated.

The client needs to fill in the following fields of `CUVIDDECODCAPS` before calling `cuvidGetDecoderCaps()`.

- ▶ **eCodecType**: Codec type (H.264, HEVC, VP9, JPEG etc.)
- ▶ **eChromaFormat**: 4:2:0, 4:4:4, etc.
- ▶ **nBitDepthMinus8**: 0 for 8-bit, 2 for 10-bit, 4 for 12-bit

When `cuvidGetDecoderCaps()` is called, the underlying driver fills up the remaining fields of `CUVIDDECODCAPS`, indicating the support for the queried capabilities, supported output formats and the maximum and minimum resolutions the hardware supports.

The following pseudo-code illustrates how to query the capabilities of NVDEC.

```
CUVIDDECODCAPS decodeCaps = {};
// set IN params for decodeCaps
decodeCaps.eCodecType = cudaVideoCodec_HEVC;//HEVC
decodeCaps.eChromaFormat = cudaVideoChromaFormat_420;//YUV 4:2:0
decodeCaps.nBitDepthMinus8 = 2;// 10 bit
result = cuvidGetDecoderCaps(&decodeCaps);
```

Returned parameters from API can be interpreted as below to validate if content can be decoded on underlying hardware:

```
// Check if content is supported
if (!decodeCaps.bIsSupported){
    NVDEC_THROW_ERROR(Codec not supported on this GPU",
    CUDA_ERROR_NOT_SUPPORTED);
}
// validate the content resolution supported on underlying hardware
if ((coded_width > decodeCaps.nMaxWidth) ||
    (coded_height > decodeCaps.nMaxHeight)){
    NVDEC_THROW_ERROR(Resolution not supported on this GPU",
    CUDA_ERROR_NOT_SUPPORTED);
}
// Max supported macroblock count CodedWidth*CodedHeight/256 must be <= nMaxMBCount
if ((coded_width>>4)*(coded_height>>4) > decodeCaps.nMaxMBCount){
    NVDEC_THROW_ERROR(MBCount not supported on this GPU",
    CUDA_ERROR_NOT_SUPPORTED);
}
```

```
}
```

In most situations, bit-depth and chroma subsampling to be used at the decoder output is same as that at the decoder input (i.e. in the content). In certain cases, however, it may be necessary to have the decoder produce output with bit-depth and chroma subsampling different from that used in the input bitstream. In general, it's always a good idea to first check if the desired output bit-depth and chroma subsampling format is supported before creating the decoder. This can be done in the following way:

```
// Check supported output format
if (decodecaps.nOutputFormatMask & (1<<cudaVideoSurfaceFormat_NV12)){
    // Decoder supports output surface format NV12
}
if (decodecaps.nOutputFormatMask & (1<<cudaVideoSurfaceFormat_P010){
    // Decoder supports output surface format P010
}
.....
```

## 4.2.2. Creating a Decoder

Before creating the decoder instance, user needs to have a valid CUDA context which will be used in the entire decoding process.

The decoder instance can be created by calling **cuvldCreateDecoder()** after filling the structure **CUVIDDECODERCREATEINFO**. The structure **CUVIDDECODERCREATEINFO** should be filled up with the following information about the stream to be decoded:

- ▶ **CodecType:** must be from **enum cudaVideoCodec**. It represents codec type of content like H.264, HEVC, VP9 etc.
- ▶ **ulWidth, ulHeight:** coded width and coded height in pixels.
- ▶ **ulMaxWidth, ulMaxHeight:** max width and max height that decoder support in case of resolution change. When there is resolution change (new resolution <= ulMaxWidth, ulMaxHeight) in video stream, app can reconfigure decoder using **cuvldReconfigureDecoder()** API instead of destroy and recreate the decoder. If ulMaxWidth or ulMaxHeight is set to 0, ulMaxWidth and ulMaxHeight are set to ulWidth and ulHeight respectively.
- ▶ **ChromaFormat:** must be from **enum cudaVideoChromaFormat**. It represents chroma format of content like 4:2:0, 4:4:4, etc.
- ▶ **bitDepthMinus8:** bit-depth minus 8 of video stream to be decoded like 0 for 8-bit, 2 for 10-bit, 4 for 12-bit.
- ▶ **ulNumDecodeSurfaces:** Referred to as decode surfaces elsewhere in this document, this is the number of surfaces that the driver will internally allocate for storing the decoded frames. Using a higher number ensures better pipelining but increases GPU memory consumption. For correct operation, minimum value is defined in **CUVIDDECODERCREATEINFO::min\_num\_decode\_surfaces** and can be obtained from first sequence callback from Nvidia parser. The NVDEC engine writes decoded data to one of these surfaces. These surfaces are not accessible by the user of NVDEC API, but the *mapping* stage, which includes decoder output format conversion, scaling, cropping etc.) use these surfaces as input surfaces.

- ▶ **ulNumOutputSurfaces:** This is the maximum number of output surfaces that the client will simultaneously *map* to decode surfaces for further processing using **cuvidMapVideoFrame()**. These surfaces have postprocessed decoded output to be used by client. The driver internally allocates the corresponding number of surfaces (referred as output surfaces in this document). Client will have access to output surfaces. Refer to section [Preparing the decoded frame for further processing](#) to understand the definition of *map*.
- ▶ **OutputFormat:** Output surface format defined as **enum cudaVideoSurfaceFormat**. This output format must be one of supported format obtained in **decodecaps.nOutputFormatMask** in **cuvidGetDecoderCaps()**. If an unsupported output format is passed, API will fail with error **CUDA\_ERROR\_NOT\_SUPPORTED**.
- ▶ **ulTargetWidth, ulTargetHeight:** This is resolution of output surfaces. For use-case which involve no scaling, these should be set to **ulWidth, ulHeight**, respectively.
- ▶ **DeinterlaceMode:** This should be set to **cudaVideoDeinterlaceMode\_Weave** or **cudaVideoDeinterlaceMode\_Bob** for progressive content and **cudaVideoDeinterlaceMode\_Adaptive** for interlaced content. **cudaVideoDeinterlaceMode\_Adaptive** yields better quality but increases memory consumption.
- ▶ **ulCreationFlags:** It is defined as **enum cudaVideoCreateFlags**. It is optional to explicitly define this flag. Driver will pick appropriate mode if not defined.
- ▶ **ulIntraDecodeOnly:** Set this flag to 1 to instruct the driver that the content being decoded contains only I/IDR frames. This helps the driver optimize memory consumption. Do not set this flag if content has non-intra frames.

The **cuvidCreateDecoder()** call fills **CUvideodecoder** with the decoder handle which should be retained till the decode session is active. The handle needs to be passed along with other NVDEC API calls.

The user can also specify the following parameters in the **CUVIDDECODERCREATEINFO** to control the final output:

- ▶ Scaling dimension
- ▶ Cropping dimension
- ▶ Dimension if the user wants to change the aspect ratio

The following code demonstrates the setup of decoder in case of scaling, cropping, or aspect ratio conversion.

```
// Scaling. Source size is 1280x960. Scale to 1920x1080.
CUresult rResult;
unsigned int uScaleW, uScaleH;
uScaleW = 1920;
uScaleH = 1080;
...
CUVIDDECODERCREATEINFO stDecodeCreateInfo;
memset(&stDecodeCreateInfo, 0, sizeof(CUVIDDECODERCREATEINFO));
... // Setup the remaining structure members
stDecodeCreateInfo.ulTargetWidth = uScaleWidth;
stDecodeCreateInfo.ulTargetHeight = uScaleHeight;
rResult = cuvidCreateDecoder(&hDecoder, &stDecodeCreateInfo);
...
```

```
// Cropping. Source size is 1280x960
CUresult rResult;
unsigned int uCropL, uCropR, uCropT, uCropB;
uCropL = 30;
uCropR = 700;
uCropT = 20;
uCropB = 500;
...
CUVIDDECODECREATEINFO stDecodeCreateInfo;
memset(&stDecodeCreateInfo, 0, sizeof(CUVIDDECODECREATEINFO));
// Setup the remaining structure members
...
stDecodeCreateInfo.display_area.left = uCropL;
stDecodeCreateInfo.display_area.right = uCropR;
stDecodeCreateInfo.display_area.top = uCropT;
stDecodeCreateInfo.display_area.bottom = uCropB;
rResult = cuvidCreateDecoder(&hDecoder, &stDecodeCreateInfo);
...

// Aspect Ratio Conversion. Source size is 1280x960(4:3). Convert to
// 16:9
CUresult rResult;
unsigned int uCropL, uCropR, uCropT, uCropB;
uDispAR_L = 0;
uDispAR_R = 1280;
uDispAR_T = 70;
uDispAR_B = 790;
...
CUVIDDECODECREATEINFO stDecodeCreateInfo;
memset(&stDecodeCreateInfo, 0, sizeof(CUVIDDECODECREATEINFO));
... // setup structure members
stDecodeCreateInfo.target_rect.left = uDispAR_L;
stDecodeCreateInfo.target_rect.right = uDispAR_R;
stDecodeCreateInfo.target_rect.top = uDispAR_T;
stDecodeCreateInfo.target_rect.bottom = uDispAR_B;
rResult = cuvidCreateDecoder(&hDecoder, &stDecodeCreateInfo);
...
```

### 4.2.3. Decoding the frame/field

After de-muxing and parsing, the client can submit the bitstream which contains a frame or field of data to hardware for decoding. To accomplish this the following steps, need to be followed:

- ▶ Fill up the **CUVIDPICPARAMS** structure.
  - ▶ The client needs to fill up the structure with parameters derived during the parsing process. **CUVIDPICPARAMS** contains a structure specific to every supported codec which should also be filled up.
- ▶ Call **cuvidDecodePicture()** and pass the decoder handle and the pointer to **CUVIDPICPARAMS**. **cuvidDecodePicture()** kicks off the decoding on NVDEC.

### 4.2.4. Preparing the decoded frame for further processing

The user needs to call **cuvidMapVideoFrame()** to get the CUDA device pointer and pitch of the output surface that holds the decoded and post-processed frame.

Please note that **cuvidDecodePicture()** instructs the NVDEC hardware engine to kick off the decoding of the frame/field. However, successful completion of

**cuvidMapVideoFrame()** indicates that the decoding process is completed and that the decoded YUV frame is converted from the format generated by NVDEC to the YUV format specified in **CUVIDDECODECREATEINFO::OutputFormat**.

**cuvidMapVideoFrame()** API takes decode surface index (**nPicIdx**) as input and *maps* it to one of available output surfaces, post-processes the decoded frame and copy to output surface and returns CUDA device pointer and associated pitch of the output surfaces.

The above operation performed by **cuvidMapVideoFrame()** is referred to as *mapping* in this document.

After the user is done with the processing on the frame, **cuvidUnmapVideoFrame()** must be called to make the output surface available for storing other decoded and post-processed frames.

If the user continuously fails to call the corresponding **cuvidUnmapVideoFrame()** after **cuvidMapVideoFrame()**, then **cuvidMapVideoFrame()** will eventually fail. At most **CUVIDDECODECREATEINFO::ulNumOutputSurfaces** frames can be mapped at a time.

**cuvidMapVideoFrame()** is a blocking call as it waits for decoding to complete. If **cuvidMapVideoFrame()** is called on same CPU thread as **cuvidDecodePicture()**, it will block **cuvidDecodePicture()** as well. In this case, the application will not be able to submit decode packets to NVDEC until mapping is complete. It can be avoided by performing the mapping operation on a CPU thread (referred as mapping thread) different from the one calling **cuvidDecodePicture()** (referred as decoding thread).

When using NVIDIA parser from NVDEC API, the application can implement a producer-consumer queue between decoding thread (as producer) and mapping thread (as consumer). The queue can contain picture indexes (or other unique identifiers) for frames being decoded. Parser can run on decoding thread. Decoding thread can add the picture index to the queue in display callback and return immediately from callback to continue decoding subsequent frames as they become available. On the other side, mapping thread will monitor the queue. If it sees the queue has non-zero length, it will dequeue the entry and call **cuvidMapVideoFrame(...)** with **nPicIdx** as the picture index. Decoding thread must ensure to not reuse the corresponding decode picture buffer for storing the decoded output until its entry is consumed and freed by mapping thread.

The following code demonstrates how to use **cuvidMapVideoFrame()** and **cuvidUnmapVideoFrame()**.

```
// MapFrame: Call cuvidMapVideoFrame and get the devptr and associated
// pitch. Copy this surface (in device memory) to host memory using
// CUDA device to host memcpy.
bool MapFrame()
{
    CUVIDPARSEDISPINFO stDispInfo;
    CUVIDPROC_PARAMS stProcParams;
    CUresult rResult;
    unsigned int cuDevPtr; int nPitch, nPicIdx;
    unsigned char* pHostPtr;
    memset(&stDispInfo, 0, sizeof(CUVIDPARSEDISPINFO));
    memset(&stProcParams, 0, sizeof(CUVIDPROC_PARAMS));
    ... // setup stProcParams if required
    // retrieve the frames from the Frame Display Queue. This Queue is
    // is populated in HandlePictureDisplay.
    if (g_pFrameQueue->dequeue(&stDispInfo))
    {
```



```

    nPicIdx = stDispInfo.picture_index;
    rResult = cuvidMapVideoFrame(&hDecoder, nPicIdx, &cuDevPtr,
                                &nPitch, &stProcParams);

    // use CUDA based Device to Host memcpy
    pHostPtr = cuMemAllocHost((void** )&pHostPtr, nPitch);
    if (pHostPtr)
    {
        rResult = cuMemcpyDtoH(pHostPtr, cuDevPtr, nPitch);
    }
    rResult = cuvidUnmapVideoFrame(&hDecoder, cuDevPtr);
}
... // Dump YUV to a file
if (pHostPtr)
{
    cuMemFreeHost(pHostPtr);
}
...
}

```

In multi-instance decoding use-case, NVDEC could be bottleneck so there wouldn't be significant benefit of calling `cuvidMapVideoFrame()` and `cuvidDecodePicture()` on different CPU threads. `cuvidDecodePicture()` will stall if wait queue on NVDEC inside driver is full. Sample applications in Video Codec SDK are using *mapping* and *decode* calls on same CPU thread, for simplicity.

## 4.2.5. Querying the decoding status

After the decoding is kicked off, `cuvidGetDecodeStatus()` can be called at any time to query the status of decoding of that frame. The underlying driver fills the status of decoding in `CUVIDGETDECODESTATUS: : *pDecodeStatus`.

The NVDEC API currently reports the following statuses:

- ▶ Decoding is in progress.
- ▶ Decoding of the frame completed successfully.
- ▶ The bitstream for the frame was corrupted and concealed by NVDEC.
- ▶ The bitstream for the frame was corrupted, however could not be concealed by NVDEC.

The API is expected to help in the scenarios where the client needs to take a further decision based on the decoding status of the frame, for e.g. whether to carry out inferencing on the frame or not.

Please note that the NVDEC can detect a limited number of errors depending on the codec. This API is supported for HEVC, H264 and JPEG on Maxwell and above generation GPUs.

## 4.2.6. Reconfiguring the decoder

Using `cuvidReconfigureDecoder()` the user can reconfigure the decoder if there is a change in the resolution and/or post processing parameters of the bitstream without having to destroy the ongoing decoder instance, and create a new one thereby saving time (and latency) in the process.

In the earlier SDKs the user had to destroy the existing decoder instance and create a new decoder instance for handling any change in decoder resolution or post processing parameters (like scaling ratio, cropping dimensions etc.).



The API can be used in scenarios where the bitstream undergoes changes in resolution, for e.g. when the encoder (on server side) changes image resolution frequently to adhere to Quality of Service(QoS) constraints.

The following steps need to be followed for using the `cuvidReconfigureDecoder()`.

1. The user needs to specify `CUVIDDECODECREATEINFO::ulMaxWidth` and `CUVIDDECODECREATEINFO::ulMaxHeight` while calling `cuvidCreateDecoder()`. The user should choose the values of `CUVIDDECODECREATEINFO::ulMaxWidth` and `CUVIDDECODECREATEINFO::ulMaxHeight` which to ensure that the resolution of the bitstream is never exceeded during the entire decoding process. Please note that the values of `CUVIDDECODECREATEINFO::ulMaxWidth` and `CUVIDDECODECREATEINFO::ulMaxHeight` cannot be changed within a session and if user wants to change the values, the decoding session should be destroyed and recreated.
2. During the process of decoding, when the user needs to change the bitstream or change postprocessing parameters, the user needs to call `cuvidReconfigureDecoder()`. This call should be ideally made from `CUVIDPARSERPARAMS::pfnSequenceCallback` when the bitstream changes. The parameters the user wants to reconfigure should be filled up in `::CUVIDRECONFIGUREDECODERINFO`. Please note, `CUVIDRECONFIGUREDECODERINFO::ulWidth` and `CUVIDRECONFIGUREDECODERINFO::ulHeight` must be equal to or smaller than `CUVIDDECODECREATEINFO::ulMaxWidth` and `CUVIDDECODECREATEINFO::ulMaxHeight` respectively or else the `cuvidReconfigureDecoder()` would fail.

The API is supported for all codecs supported by NVDEC API.

## 4.2.7. Destroying the decoder

The user needs to call `cuvidDestroyDecoder()` to destroy the decoder session and free up all the allocated decoder resources.

## 4.3. Writing an Efficient Decode Application

The NVDEC engine on NVIDIA GPUs is a dedicated hardware block, which decodes the input video bitstream in supported formats. A typical video decode application broadly consists of the following stages:

1. De-Muxing
2. Video bitstream parsing and decoding
3. Preparing the frame for further processing

Of these, de-muxing and parsing are not hardware accelerated and therefore outside the scope of this document. The de-muxing can be performed using third party components such as FFmpeg, which provides support for many multiplexed video formats. The sample applications included in the SDK demonstrate de-muxing using FFmpeg.

Similarly, post-decode or video post-processing (such as scaling, color space conversion, noise reduction, color enhancement etc.) can be effectively performed using user-defined CUDA kernels.

The post-processed frames can then be sent to the display engine for displaying on the screen, if required. Note that this operation is outside the scope of NVDEC APIs.

An optimized implementation should use independent threads for de-muxing, parsing, bitstream decode and processing etc. as explained below:

1. De-muxing: This thread demultiplexes the media file and makes the raw bit-stream available for parser to consume.
2. Parsing and decoding: This thread does the parsing of the bitstream and kicks off decoding by calling `cuvidDecodePicture()`.
3. Mapping and making the frame available for further processing: This thread checks if there are any decoded frames available. If yes, then it should call `cuvidMapVideoFrame()` to get the CUDA device pointer and pitch of the frame. The frame can then be used for further processing.

The NVDEC driver internally maintains a queue of 4 frames for efficient pipelining of operations. Please note that this pipeline does not imply any decoding delay for decoding. The decoding starts as soon as the first frame is queued, but the application can continue queuing up input frames so long as space is available without stalling. Typically, by the time application has queued 2-3 frames, decoding of the first frame is complete and the pipeline continues. This pipeline ensures that the hardware decoder is utilized to the maximum extent possible.

For performance intensive and low latency video codec applications, ensure the PCIe link width is set to the maximum available value. PCIe link width currently configured can be obtained by running command 'nvidia-smi -q'. PCIe link width can be configured in the system's BIOS settings.

In the use cases where there is frequent change of decode resolution and/or post processing parameters, it is recommended to use `cuvidReconfigureDecoder()` instead of destroying the existing decoder instance and recreating a new one.

The following steps should be followed for optimizing video memory usage:

1. Make `CUVIDDECODERCREATEINFO::ulNumDecodeSurfaces = CUVIDDECODERCREATEINFO::min_num_decode_surfaces`. This will ensure that the underlying driver allocates minimum number of decode surfaces to correctly decode the sequence. In case there is reduction in decoder performance, clients can slightly increase `CUVIDDECODERCREATEINFO::ulNumDecodeSurfaces`. It is therefore recommended to choose the optimal value of `CUVIDDECODERCREATEINFO::ulNumDecodeSurfaces` to ensure right balance between decoder throughput and memory consumption.
2. `CUVIDDECODERCREATEINFO::ulNumOutputSurfaces` should be decided optimally after due experimentation for balancing decoder throughput and memory consumption.
3. `CUVIDDECODERCREATEINFO::DeinterlaceMode` should be set to `"cudaVideoDeinterlaceMode::cudaVideoDeinterlaceMode_Weave"` or `"cudaVideoDeinterlaceMode::cudaVideoDeinterlaceMode_Bob"`. For interlaced contents, choosing

`cudaVideoDeinterlaceMode::cudaVideoDeinterlaceMode_Adaptive` results to higher quality but increases memory consumption. Using `cudaVideoDeinterlaceMode::cudaVideoDeinterlaceMode_Weave` or `cudaVideoDeinterlaceMode::cudaVideoDeinterlaceMode_Bob` results to minimum memory consumption though it may result in lesser video quality. In case “`CUVIDDECODERCREATEINFO::DeinterlaceMode`” is not specified by the client, the underlying display driver sets it to “`cudaVideoDeinterlaceMode::cudaVideoDeinterlaceMode_Adaptive`” which results to higher memory consumption. Hence it is strongly recommended to choose the right value of `CUVIDDECODERCREATEINFO::DeinterlaceMode` depending on the requirement.

4. While decoding multiple streams it is recommended to allocate minimum number of CUDA contexts and share it across sessions. This saves the memory overhead associated with the CUDA context creation.
5. `CUVIDDECODERCREATEINFO::ulIntraDecodeOnly` should be set to 1 if it is known beforehand that the sequence contains Intra frames only. This feature is supported only for HEVC, H.264 and VP9. However, decoding might fail if the flag is enabled in case of supported codecs for regular bit streams having P and/or B frames.

The sample applications included with the Video Codec SDK are written to demonstrate the functionality of various APIs, but they may not be fully optimized. Hence programmers are strongly encouraged to ensure that their application is well-designed, with various stages in the decode-postprocess-display pipeline structured in an efficient manner to achieve desired performance and memory consumption.

ALL NVIDIA DESIGN SPECIFICATIONS, REFERENCE BOARDS, FILES, DRAWINGS, DIAGNOSTICS, LISTS, AND OTHER DOCUMENTS (TOGETHER AND SEPARATELY, "MATERIALS") ARE BEING PROVIDED "AS IS." NVIDIA MAKES NO WARRANTIES, EXPRESSED, IMPLIED, STATUTORY, OR OTHERWISE WITH RESPECT TO THE MATERIALS, AND EXPRESSLY DISCLAIMS ALL IMPLIED WARRANTIES OF NONINFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE.

Information furnished is believed to be accurate and reliable. However, NVIDIA Corporation assumes no responsibility for the consequences of use of such information or for any infringement of patents or other rights of third parties that may result from its use. No license is granted by implication of otherwise under any patent rights of NVIDIA Corporation. Specifications mentioned in this publication are subject to change without notice. This publication supersedes and replaces all other information previously supplied. NVIDIA Corporation products are not authorized as critical components in life support devices or systems without express written approval of NVIDIA Corporation.

HDMI, the HDMI logo, and High-Definition Multimedia Interface are trademarks or registered trademarks of HDMI Licensing LLC.

OpenCL is a trademark of Apple Inc. used under license to the Khronos Group Inc.

NVIDIA and the NVIDIA logo are trademarks or registered trademarks of NVIDIA Corporation in the U.S. and other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

© 2010-2020 NVIDIA Corporation. All rights reserved.